



Paper Summaries

Some Joke About VAEs

Written @ Corti

Authors:
Magnus Berg Sletfjording
{ms}@corti.ai
May 15, 2021



Abstract

This was written for me to understand papers in my thesis better. Don't be alarmed if you don't understand it 100%, I probably don't either.



Contents

0	Auto Encoding Variational Bayes	1
0.1	Method and setup	1



Table 1: Overview of the different definitions in the VAE

Distribution	Equation
Generative model	$p(\mathbf{x}) = p_{\theta}(\mathbf{x} \mathbf{z})p_{\theta}(\mathbf{z})$
Prior over \mathbf{z}	$p_{\theta}(\mathbf{z})$
Likelihood over \mathbf{x}	$p_{\theta}(\mathbf{x} \mathbf{z})$
Posterior over \mathbf{z}	$p_{\theta}(\mathbf{z} \mathbf{x})$
Posterior approximation over \mathbf{z} $q_{\psi}(\mathbf{z} \mathbf{x})$	

0 Auto Encoding Variational Bayes

Authors: Diederik P. Kingma, Max Welling [?]

Main contributions of this paper is as follows:

1. The reparametrization trick for SGD methods
2. The first "vanilla" VAE

0.1 Method and setup

Assume a directed graphical model able to do both inference and running a pure generative process.

We define the problem as the following: Given a generative model $p_{\theta}(x)$, we want to find the optimal parameters for θ .

$$p(\mathbf{x}) = p_{\theta}(\mathbf{x}|\mathbf{z})p_{\theta}(\mathbf{z}) \tag{1}$$

References

- [1] J. Chung, S. Ahn, and Y. Bengio. Hierarchical Multiscale Recurrent Neural Networks. arXiv:1609.01704 [cs], Mar. 2017. arXiv: 1609.01704.
- [2] T. Kim, S. Ahn, and Y. Bengio. Variational Temporal Abstraction. arXiv:1910.00775 [cs, stat], Oct. 2019. arXiv: 1910.00775.
- [3] J. Koutník, K. Greff, F. Gomez, and J. Schmidhuber. A Clockwork RNN. arXiv:1402.3511 [cs], Feb. 2014. arXiv: 1402.3511.
- [4] A. v. d. Oord, S. Dieleman, H. Zen, K. Simonyan, O. Vinyals, A. Graves, N. Kalchbrenner, A. Senior, and K. Kavukcuoglu. WaveNet: A Generative Model for Raw Audio. arXiv:1609.03499 [cs], Sept. 2016. arXiv: 1609.03499.
- [5] A. v. d. Oord, O. Vinyals, and K. Kavukcuoglu. Neural Discrete Representation Learning. arXiv:1711.00937 [cs], May 2018. arXiv: 1711.00937.
- [6] A. Radford, L. Metz, and S. Chintala. Unsupervised Representation Learning with Deep Convolutional Generative Adversarial Networks. arXiv:1511.06434 [cs], Jan. 2016. arXiv: 1511.06434.
- [7] V. Saxena, J. Ba, and D. Hafner. Clockwork Variational Autoencoders. arXiv:2102.09532 [cs], Feb. 2021. arXiv: 2102.09532.
- [8] D. Shen, A. Celikyilmaz, Y. Zhang, L. Chen, X. Wang, J. Gao, and L. Carin. Towards Generating Long and Coherent Text with Multi-Level Latent Variable Models. In Proceedings of the 57th Annual Meeting of the Association for Computational Linguistics, pages 2079–2089, Florence, Italy, July 2019. Association for Computational Linguistics.